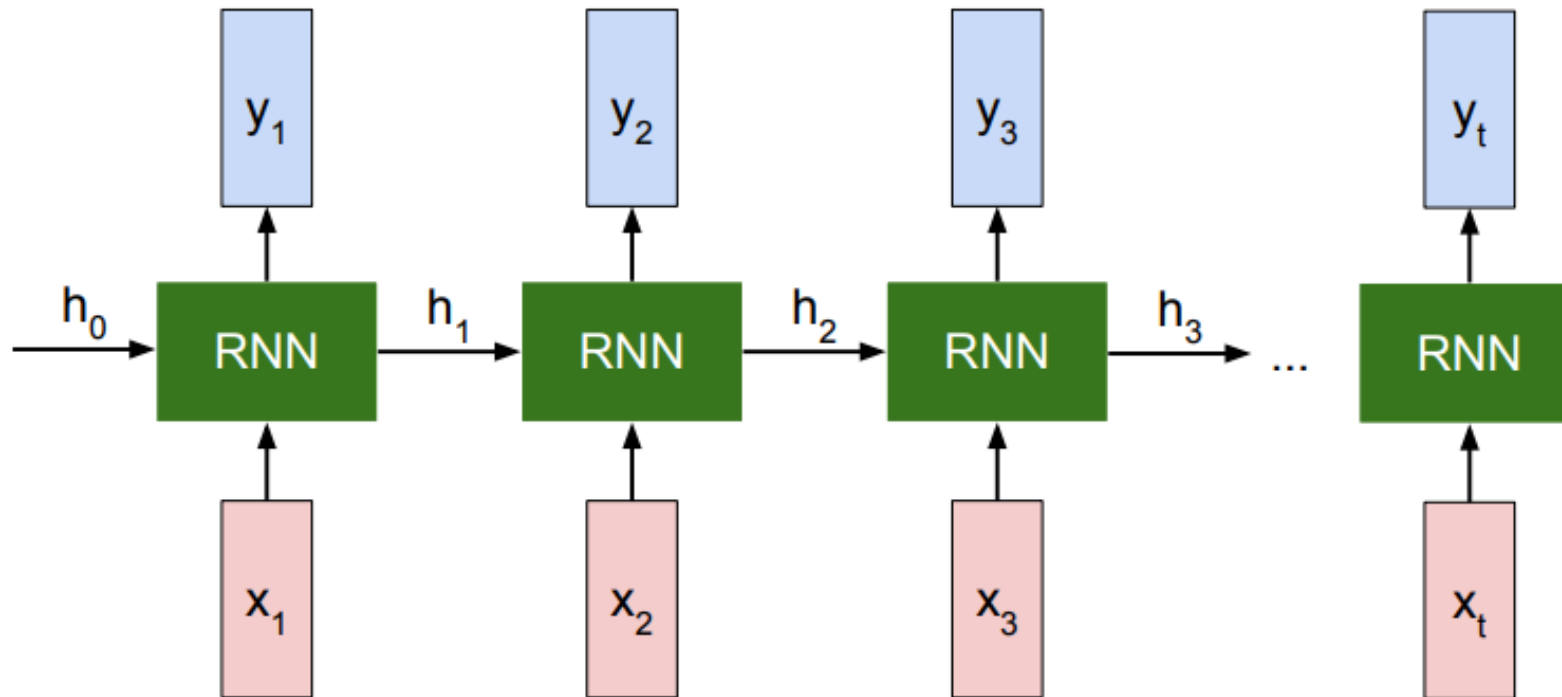


# Task-Optimized RNN in PyTorch

•

Choice Phase: Select eel with larger influence

# RNN



Task

Competency Task Visualization

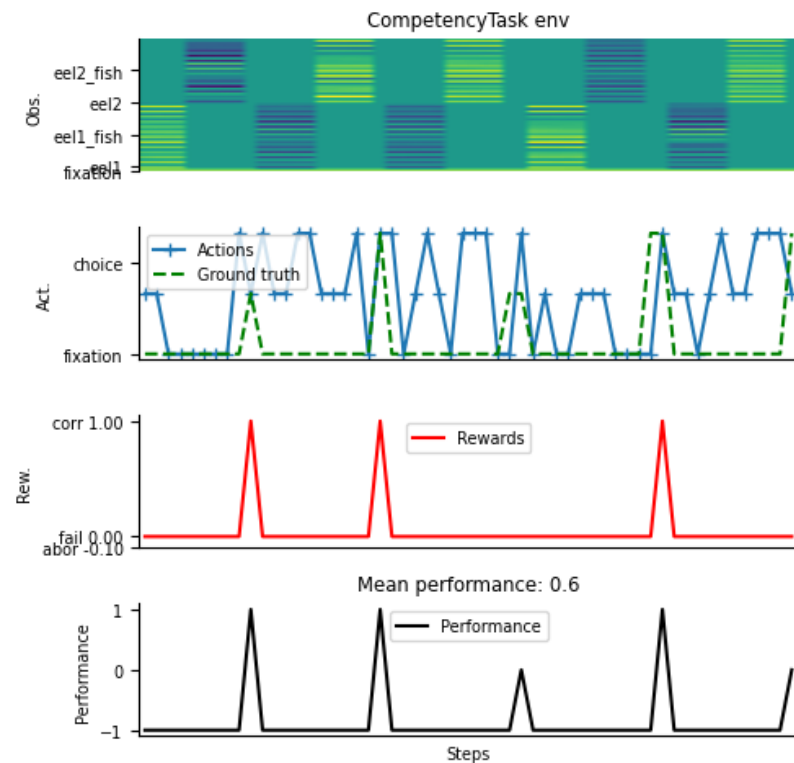


**P: Toggle Potential Field, SPACE: Pause, R: Reset, ESC: Quit**



# Breaking Down the CompetencyTask

## Example Network



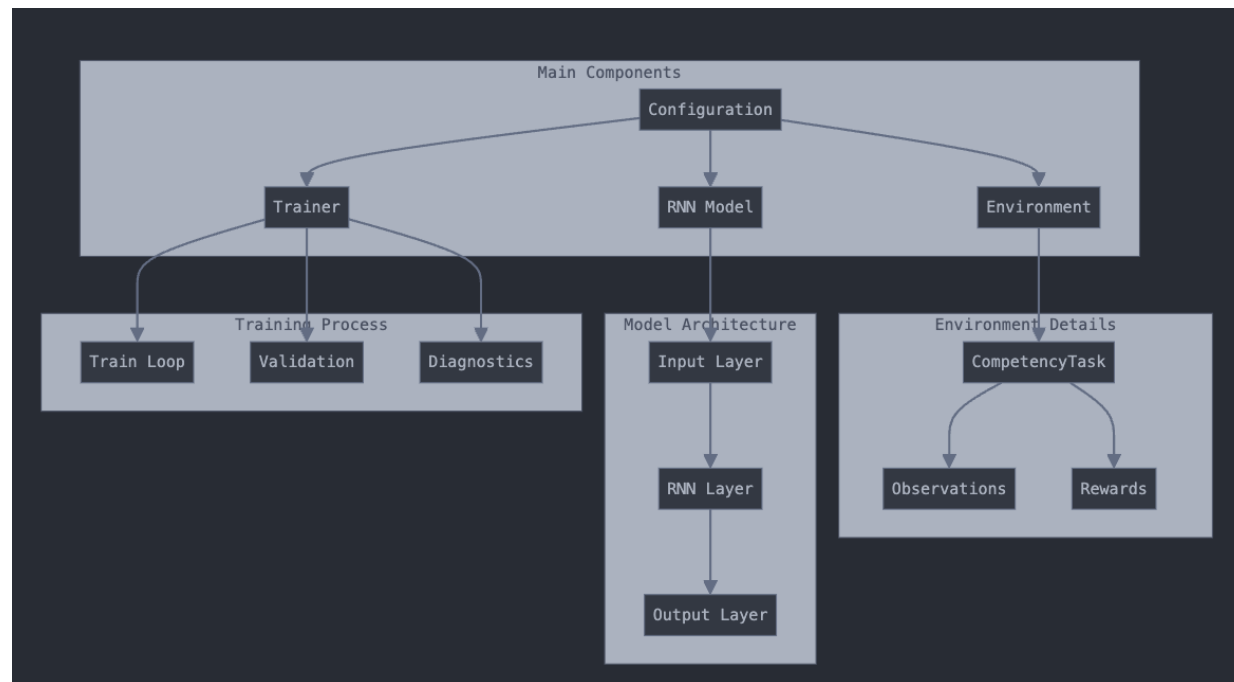
# Code

Main

# Project Structure

▼	📁	Tutorial	●
▼	📁	models	●
>	📁	__pycache__	
	🐍	__init__.py	
	🐍	students_rnn.py	2
	🐍	competency_task_students.py	6
	📄	config.yaml	
	🐍	main.py	5
	🐍	trainer.py	4

# Main

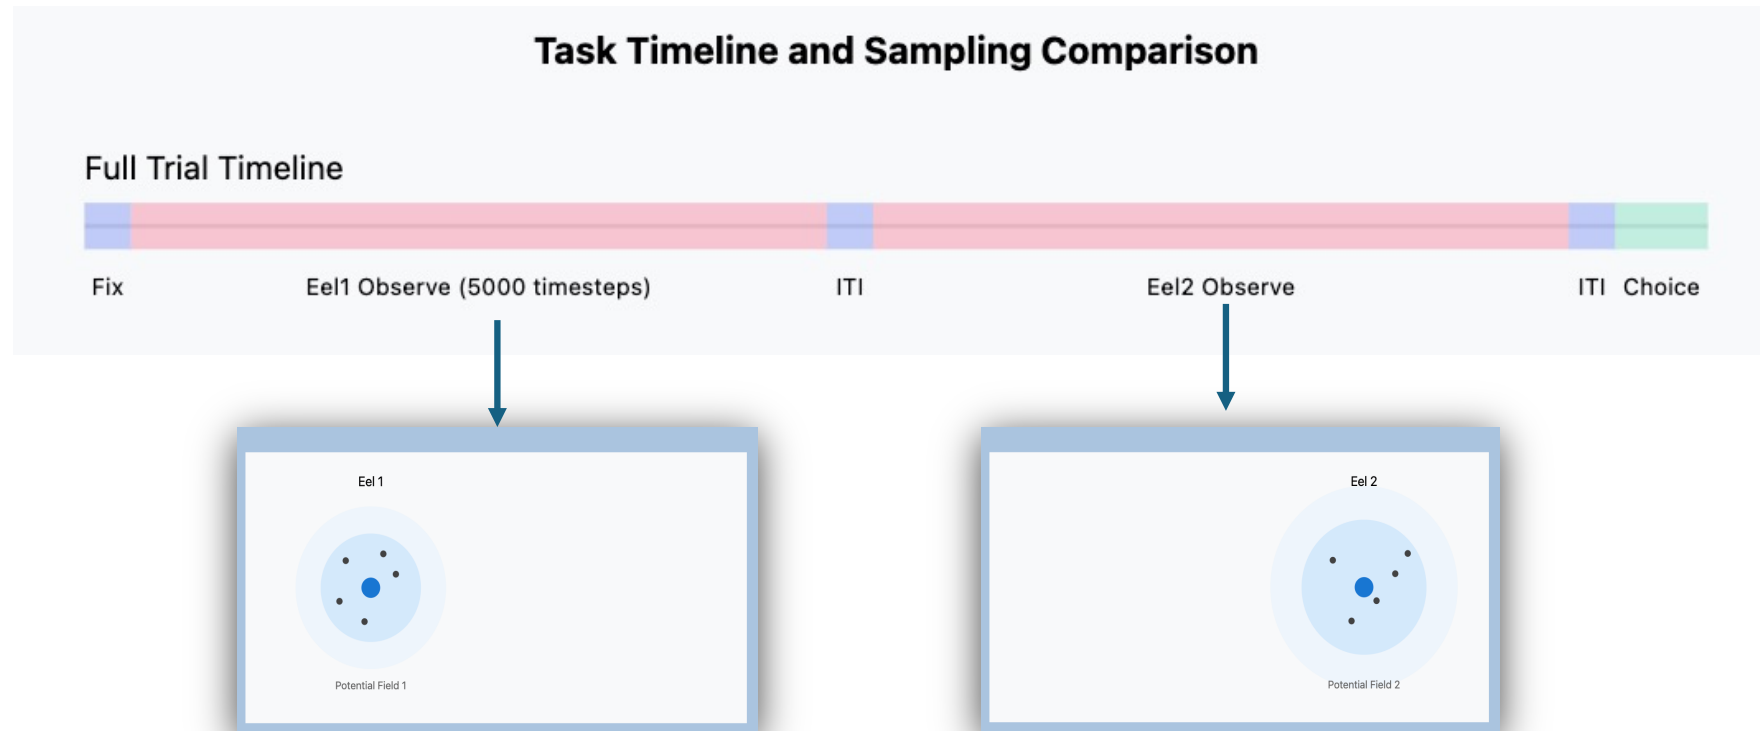




# Code

Environment


# Environment



# Environment

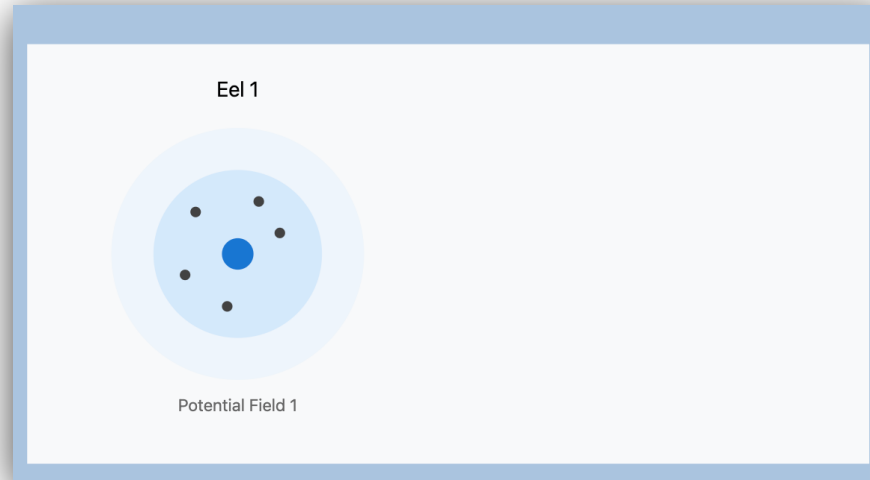
## Task

- **What is the Task?**
  - The agent must:
    - Observe two eels and their influence on nearby fish
    - Determine which eel has a stronger influence
    - Make a correct choice to receive a reward
- **Trial Structure**
  - 1. Fixation Phase**
    1. Nothing shown
  - 2. First Observation**
    1. Watch first eel and its fish
  - 3. Second Observation**
    1. Watch second eel and its fish
  - 4. Decision Phase**
    1. Choose which eel has stronger influence
    2. Receive reward if correct




# Environment

## Timepoint Example



```
# During eel1_observe:  
obs = [1, x1, y1, fish1_x1, fish1_y1, ..., fish12_x1, fish12_y1, 0, 0, ..., 0]  
  
# During eel2_observe:  
obs = [1, 0, 0, ..., 0, x2, y2, fish1_x2, fish1_y2, ..., fish12_x2, fish12_y2]
```



Total Dimensions: 53

# Code

Configuration

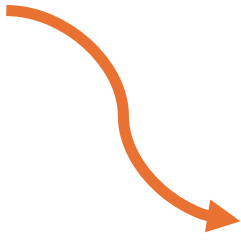
# Configuration (config.yaml)

## Model

```
# config.yaml

# Model Configuration
model:
  hidden_size: 64
  dtype: float32
```

- **hidden\_size: Controls model capacity**
  - **Larger** = more complex patterns but slower training
  - **Smaller** = faster training but might miss subtle patterns

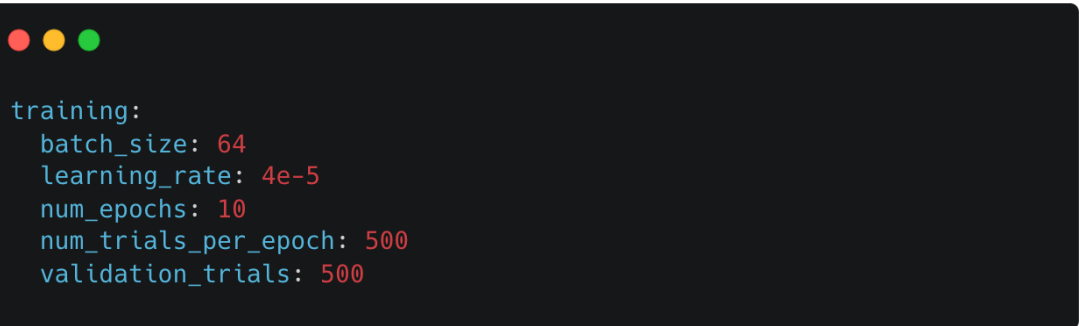


```
# main.py

# Initialize model
model = RNNNet(
    input_size=env.observation_space.shape[0],
    hidden_size=config['model']['hidden_size'],
    output_size=env.action_space.n,
).to(device)
```

# Configuration (config.yaml)

## Trainer



```
training:
  batch_size: 64
  learning_rate: 4e-5
  num_epochs: 10
  num_trials_per_epoch: 500
  validation_trials: 500
```

- **batch\_size**: Number of trials processed together
    - **Larger** = more stable gradients but more memory
    - **Smaller** = less stable but faster updates
  - **learning\_rate**: Controls step size in optimization
    - **Too high** = unstable training
    - **Too low** = slow convergence
- num\_epochs**: Total training iterations
- validation\_frequency**: How often to check performance

# Definitions

## Gradients

Gradients drive optimization by:

1. Quantifying how each parameter affects the loss
2. Determining parameter updates during backpropagation
3. Enabling efficient weight optimization through gradient descent

```
● ● ●  
# Configuration parameters affecting gradients  
training:  
  learning_rate: 4e-5 # Step size for updates  
  optimizer:  
    type: 'adamw' # Adaptive optimization  
    weight_decay: 1e-4 # L2 regularization
```

### •Gradient Magnitude

- Large gradients → Risk of unstable training
- Very small gradients → Slow convergence or vanishing gradient problem
- Clipping prevents extreme updates while preserving direction

### •Optimization Strategy

- Learning rate:** scales gradient-based updates
- AdamW:** Adaptive optimizer with momentum
- Weight decay:** Prevents weights from growing too large



# Configuration(config.yaml)

## Task

```
# Task Configuration
task:
  dt: 10 # Keep 50hz sampling rate
  timing:
    fixation: 100
    eel1_observe: 5000 # Doubled observation time
    iti1: 100
    eel2_observe: 5000 # Doubled observation time
    iti2: 100
    choice: 100

# Environment Parameters
environment:
  eel:
    speed: 0.02
    size: 0.1
    wiggle: 0.01

  fish:
    count: 12
    size: 0.05
    speed:
      slow: 0.002
      fast: 0.05
```



### •dt = 20ms

- For fixation period (100 timesteps):
  - 2000ms total duration
  - 100 observations (2000/20)

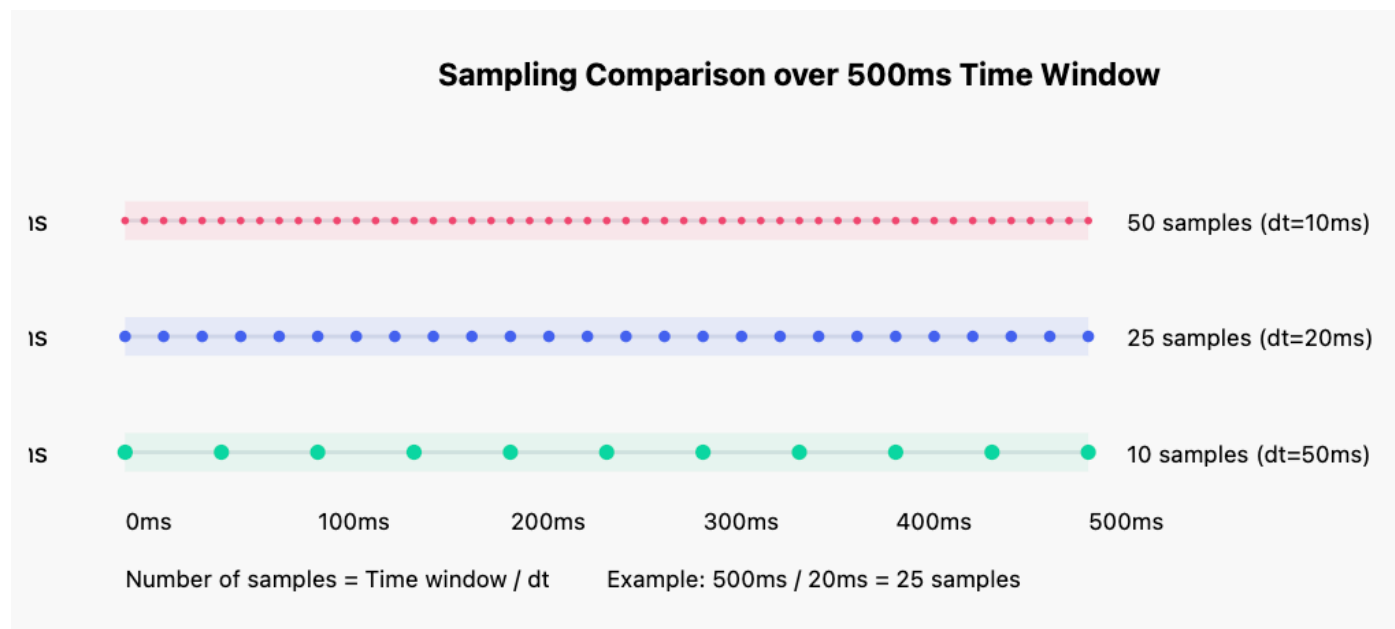
### •dt = 10ms

- For fixation period (100 timesteps):
  - 1000ms total duration
  - 200 observations (1000/10)

# Configuration


DT delta time

- Time step used to advance the environment's state in each simulation step.
- Specifies the amount of that passes between updates in the environment, defining the resolution of time.



# Code

Trainer



```
def __init__(self, model, env, config):
    self.model = model          # Neural network model
    self.env = env              # Task environment
    self.config = config        # Configuration parameters

    # Training metrics
    self.training_losses = []
    self.training_accuracies = []
    self.validation_accuracies = []

    # Key parameters from config
    self.num_epochs = config['training']['num_epochs']
    self.batch_size = config['training']['batch_size']
    self.num_trials_per_epoch = config['training']['num_trials_per_epoch']
```

# Trainer

## train



```
def train(self):
    for epoch in range(self.num_epochs):
        # 1. Process batches
        for batch_idx in range(num_batches):
            # Get batch of trials
            inputs_batch, labels_batch = self.get_batch()

            # Train on this batch
            loss = self.train_step(inputs_batch, labels_batch)

            # Track performance
            accuracy = self.evaluate_batch()

        # 2. Run validation periodically
        if (epoch + 1) % self.val_frequency == 0:
            val_accuracy = self.validate()

        # 3. Run diagnostics periodically
        if (epoch + 1) % 10 == 0:
            self.run_diagnostics()
```

- Epoch management
- Batch processing
- Periodic validation
- Progress tracking

# Trainer

## train\_step

```
def train_step(self, inputs, labels):  
    """Single training step with improvements."""  
    # Zero gradients  
    self.optimizer.zero_grad()  
  
    # Normalize & forward pass  
    inputs = (inputs - inputs.mean()) / (inputs.std() + 1e-8)  
    outputs, _ = self.model(inputs)  
  
    # Loss & backprop  
    loss = self.criterion(outputs, labels)  
    loss.backward()  
  
    self.optimizer.step()  
  
    return loss.item()
```

- Forward: Get prediction, loss  
Backward: Compute gradients  
Update: Improve weights

# Trainer

## train\_step

```
def train_step(self, inputs, labels):  
    # Zero gradients  
    self.optimizer.zero_grad()  
  
    # Forward pass  
    outputs, _ = self.model(inputs)  
  
    # Calculate loss  
    loss = self.criterion(outputs, labels)  
  
    # Backward pass with gradient clipping  
    loss.backward()  
    torch.nn.utils.clip_grad_norm_(  
        self.model.parameters(),  
        self.config['training']['gradient_clip']  
    )  
  
    # Update weights  
    self.optimizer.step()
```

# Trainer

## Validation

```
def validate(self):
    self.model.eval()
    correct = 0
    total = 0

    with torch.no_grad():
        for _ in range(self.validation_trials):
            # Run trial
            trial_info = self.env.new_trial()

            # Get prediction
            inputs = self.process_inputs(self.env.ob)
            outputs, _ = self.model(inputs)

            # Check accuracy
            prediction = self.get_prediction(outputs)
            if prediction == trial_info['ground_truth']:
                correct += 1
            total += 1

    return correct / total
```

- Evaluate generalization
- Track learning progress
- Prevent overfitting

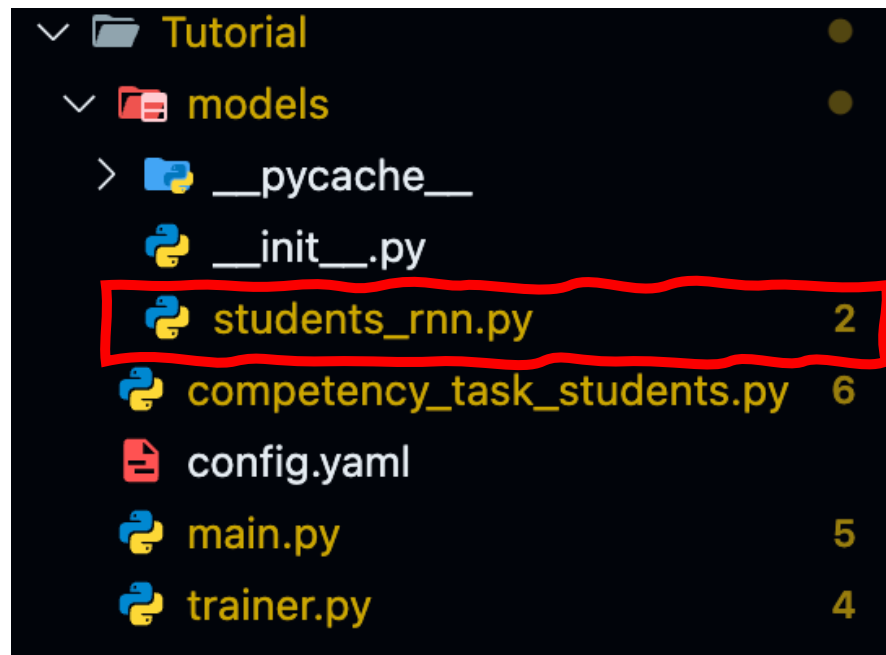


# Code

Model

# Assignment

Model



A screenshot of a file explorer interface with a dark background. The file structure is as follows:

- ▼ Tutorial
  - ▼ models
    - > \_\_pycache\_\_
    - \_\_init\_\_.py
    - students\_rnn.py 2** (highlighted with a red box)
    - competency\_task\_students.py 6
    - config.yaml
    - main.py 5
    - trainer.py 4

File/Folder	Count
students_rnn.py	2
competency_task_students.py	6
main.py	5
trainer.py	4